




Beyond Low-rank Decomposition: A Shortcut Approach for Efficient On-Device Learning

Le-Trung Nguyen Aël Quélenec Van-Tam Nguyen Enzo Tartaglione

LTCI, Télécom Paris, Institut Polytechnique de Paris

{name.surname}@telecom-paris.fr



How to save up to **120.09× activation memory** and **1.86× computational cost** during training?


The Memory Bottleneck of Backpropagation

Considering a convolutional neural network, during the backward pass at the i^{th} layer, the following two values must be calculated:

$$\frac{\partial \mathcal{L}}{\partial \mathcal{W}_i} = \frac{\partial \mathcal{L}}{\partial \mathcal{A}_{i+1}} \cdot \frac{\partial \mathcal{A}_{i+1}}{\partial \mathcal{W}_i} = \text{conv}\left(\mathcal{A}_i, \frac{\partial \mathcal{L}}{\partial \mathcal{A}_{i+1}}\right)$$


$$\frac{\partial \mathcal{L}}{\partial \mathcal{A}_i} = \frac{\partial \mathcal{L}}{\partial \mathcal{A}_{i+1}} \cdot \frac{\partial \mathcal{A}_{i+1}}{\partial \mathcal{A}_i} = \text{convfull}\left[\frac{\partial \mathcal{L}}{\partial \mathcal{A}_{i+1}}, \text{rot}(\mathcal{W}_i)\right]$$

Storing \mathcal{A}_i and \mathcal{W}_i is the main cause of memory occupancy during backpropagation. In this work, we propose decomposing \mathcal{A}_i during the forward pass using **Activation Subspace Iteration (ASI)**. \mathcal{W}_i is untouched to avoid error propagation through $\frac{\partial \mathcal{L}}{\partial \mathcal{A}_i}$.




Method

ASI in Backpropagation




ASI flowchart




Expected results

- Consider activation maps of the i^{th} and the $(i+1)^{th}$ layers as $\mathcal{A}_i \in \mathbb{R}^{B \times C_i \times H_i \times W_i}$ and $\mathcal{A}'_i \in \mathbb{R}^{B \times C'_i \times H'_i \times W'_i}$, respectively.
- Let $\mathcal{D}_{i,m} = \{B, C_i, H_i, W_i\}$ and $\mathcal{D}'_{i,m} = \{B, C'_i, H'_i, W'_i\}$.
- The target rank for compression is $r \in \mathbb{R}^{N \times 4}$, where N is total number of layers and 4 is number of tensor mode of the activation maps.




Results

Fine-tuning different number of layers of MCUNet pretrained on ImageNet with CIFAR-10 as downstream dataset



- ASI has similar performance to HOSVD in terms of memory compression, but it is significantly more efficient.
- For the same accuracy, ASI saves 120.09× more memory and is 1.86× cheaper than vanilla training.

Processing time



- The forward processing time of ASI is approximately the same as that of vanilla training, while the backward pass is up to an order of magnitude faster.
- Overall, ASI is on average 1.5× faster than vanilla training.


Acknowledgement: Part of this work was funded by HiPARIS Center on Data Analytics and Artificial Intelligence, by the European Union's Horizon Europe Research and Innovation Programme under grant agreement No. 101120237 (ELIAS), by the European Union's HORIZON Research and Innovation Programme under grant agreement No 101120657, project ENFIELD (European Lighthouse to Manifest Trustworthy and Green AI) and by French National Research Agency (ANR-22-PEFT-0003 and ANR-22-PEFT-0007) as part of France 2030, the NF-NAI project and NF-FITNESS project.

Experiment on ImageNet classification with different models

Method	MobileNetV2			Method	ResNet18			
	#Layers	Acc ↑	Mem (MB) ↓		#Layers	Acc ↑	Mem (MB) ↓	GFLOPs ↓
Vanilla training	All	74.0	1651.84	830.82	Vanilla training	72.8	532.88	291.79
	2	62.6	15.31	4.50		69.9	12.25	29.60
Gradient filtering R2	2	62.6	5.00	4.50	Gradient filtering R2	68.7	4.00	29.60
	4	65.2	9.38	57.50		69.3	7.00	47.70
HOSVD ($\epsilon = 0.8$)	2	61.1	0.15	3049.71	HOSVD ($\epsilon = 0.8$)	69.2	0.97	1581.42
	4	63.9	0.73	5895.31		70.5	2.89	4048.56
ASI	2	60.3	0.13	2.81	ASI	68.9	0.93	19.04
	4	64.0	0.71	31.54		70.6	2.63	33.14

Method	MCUNet			Method	ResNet34			
	#Layers	Acc ↑	Mem (MB) ↓		#Layers	Acc ↑	Mem (MB) ↓	GFLOPs ↓
Vanilla training	All	67.4	632.98	248.84	Vanilla training	75.6	839.04	528.55
	2	62.1	13.78	19.31		69.6	12.25	29.60
Gradient filtering R2	4	64.7	19.52	19.88	Gradient filtering R2	72.2	24.50	59.19
	2	61.8	4.50	19.31		68.8	4.00	29.60
HOSVD ($\epsilon = 0.8$)	4	64.4	6.38	19.89	HOSVD ($\epsilon = 0.8$)	70.9	8.00	59.21
	2	61.7	0.48	1988.05		68.7	0.30	1579.64
ASI	4	63.9	0.88	2457.59	ASI	71.1	1.11	3160.46
	2	61.7	0.38	11.01		68.3	0.25	16.44
	4	63.5	0.83	11.93		71.1	1.09	35.31

For the same #Layers, ASI saves significantly more memory and GLOPS than the others.



Paper

Code